SLURM

GPU Server Use Documentation

Table of Contents

GPU Server Use Documentation

First Things to do: (Not necessary)

1. `module load anaconda/3`
2. `conda init`
3. `source ~/.bashrc`

Now close the terminal and open a new one. You will get -

```
(base) dborman@headnode:~$
```

something like this.

Using Gpu Node:

Make script like this with `.sh` extension on the project folder(you need to give proper path for the python/r script)-

```
#!/bin/bash
#SBATCH --partition=gpu            # partition name
#SBATCH --nodes=1                     # number of nodes you want
#SBATCH --ntasks=1                     # submitting number of batch tasks (1
#SBATCH --cpus-per-task=5             # number of CPU cores per task
#SBATCH --gres=gpu:1                  # number of gpu you want
#SBATCH --mem=10GB                    # amount of RAM
#SBATCH --time=120:00:00              # maximum runtime (adjust if needed)
#SBATCH --job-name=test-slurm-gpu     # job name
#SBATCH --output=output-test-gpu.log  # output log file
#SBATCH --mail-type=BEGIN,END,FAIL    # email on job begin, end, and failure
#SBATCH --mail-user=youremail@example.com # your email address

# very important configs
module load anaconda/3  #load anaconda
conda init              # initialize it
source $(conda info --base)/etc/profile.d/conda.sh  # load conda on profile
conda activate ml # your desired environment


# Print start time
echo "Job started at: $(date)"

# Run the Python script
python test.py

# Print end time
echo "Job finished at: $(date)"
```

Then subumit job: `sbatch script_file.sh`

Now your job has been submitted.

To view your job status use `squeue`

To cancel a job use `scancel JOB_ID` (use squeue to see it easily).